

最近在部署FLUME监控日志tomcat日志文件到hadoop的hdfs，发现了一个奇怪的问题：flume使用exec方式监控tomcat单个日志文件，如果tomcat的日志文件按照天进行滚动的话，比如今天的日志为catalina.2017-08-05.out 那么过了今天，tomcat的日志文件变为了catalina.2017-08-06.out，但是flume监控配置并不能感知到tomcat日志的切换，监控的还是上一天的日志。

原来的flume的配置如下：

```
agent1.sources.s2.command = tail -n +0 -F
"/home/gome_guest/10.58.61.83/cashier-service_02/logs/catalina.`date +%Y-%m-%d`.out"
```

更改后的配置为：

```
agent1.sources.s2.command = locktail_rotate.sh
/home/gome_guest/10.58.61.83/cashier-
service_02/logs/catalina.DATE_ROTATE.out 'date +"%Y-%m-%d"'
```

其中 locktail_rotate.sh 参见

https://github.com/ypenglyn/locktail/blob/master/locktail_rotate.sh

```
#!/bin/bash
```

```
#
```

```
# AUTHOR: Yanpeng Lin
```

```
# DATE: Mar 30 2014
```

```
# DESC: lock a rotating file(filename rotated by date) and tail
```

```
#
```

```
PID=$( mktemp )
```

```
while true;
```

```
do
```

```
    CURRENT_PATTERN=$( eval "$2"
```

```
    CURRENT_PATH="$1"
```

```
    CURRENT_TARGET=$( echo ${CURRENT_PATH} | sed
"s/DATE_ROTATE/$CURRENT_PATTERN/g" )
```

```

# echo $CURRENT_TARGET
if [ -e ${CURRENT_TARGET} ]; then
    IO=`stat -c %i ${CURRENT_TARGET}`
    tail -f ${CURRENT_TARGET} 2> /dev/null & echo $! > ${PID};
fi

# as long as the file exists and the inode number did not change
while [[ -e ${CURRENT_TARGET} ]] && [[ ${IO} = `stat -c %i
${CURRENT_TARGET}` ]]
do
    CURRENT_PATTERN=$( eval "$2" )
    CURRENT_PATH="$1"
    CURRENT_TARGET=$( echo ${CURRENT_PATH} | sed
"s/DATE_ROTATE/${CURRENT_PATTERN}/g" )
    #echo $CURRENT_TARGET
    sleep 0.5
done
# echo "[ ! -z $PID ] && kill `cat $PID` 2> /dev/null && echo > $PID"
if [ ! -z ${PID} ]; then
    kill `cat ${PID}` 2> /dev/null && echo > ${PID}
fi
sleep 0.5
done 2> /dev/null
rm -rf ${PID}

```